

Himanshu Manish Jain

+1 (303) 641-6694 | himanshumjain15@gmail.com | [LinkedIn](#) | [Portfolio](#) | Boulder, CO | Open to Relocate

EDUCATION

University of Colorado Boulder

Master of Science in Data Science

- Relevant Courses: Statistical Inference, Data Mining, Big Data Architecture

Manipal Institute of Technology

Bachelor of Science in Electronics and Communication

Boulder, CO

Aug. 2025 – May 2027

India

Aug. 2021 – May 2025

SKILLS

Programming & Tools: Python, SQL, PostgreSQL, Docker, FastAPI, AWS, Tableau, Power BI

Data & ML: Pandas, NumPy, scikit-learn, PyTorch, NLP, LLM APIs

Statistical Methods: Statistical Modeling, Hypothesis Testing, A/B Testing, Outlier Analysis, EDA

Business Skills: Cross-Functional Collaboration, Stakeholder Management, Technical Documentation

EXPERIENCE

Summer Atlantic Capital

Data Science Intern

Newark, DE, USA

June 2026 – August 2026

- Built a **Python/SQL** platform for self-service compensation lookups, using precomputation and two-tier caching over **790K** market records to cut lookup time from **10–20 minutes** to seconds
- Surfaced org-wide pay-equity gaps via **Tableau** dashboards backed by a **MAD/modified-Z** outlier detection and privacy-suppression engine
- Integrated **Claude API**-driven entity resolution with human-in-the-loop review, strengthening match reliability across internal and external market-record sources
- Restored data integrity for compensation benchmarking by fixing two silent pipeline defects, recovering **~1,600** dropped market records

Ernst and Young (EY)

Data Science Intern

Bengaluru, KA, India

Jan. 2025 – July 2025

- Validated a **22%** engagement lift by running **A/B tests** on **100K+** platform users to evaluate an auto call routing feature, using a **two-sample t-test** to establish significance before scaling
- Increased adoption by **30%** on Tele-MANAS, a **national** mental health platform, by turning **10M+** records into actionable insights via **SQL** to surface regional drop-off patterns
- Partnered with medical, engineering, and government stakeholders to translate drop-off insights into **cross-functional** rollouts across Tele-MANAS
- Reduced strategic decision cycles by **25%** for **10+** regional teams by designing and deploying **Power BI** dashboards to track performance KPIs

LG Electronics

Data Analyst Intern

Noida, UP, India

May 2024 – June 2024

- Identified a **10%** cost-saving opportunity by applying **z-score outlier analysis** to flag suppliers with anomalous raw material costs, enabling data-driven negotiations
- Accelerated contract approvals by **15%** by designing a scalable **Python + SQL** pipeline to integrate and process supplier **KPI** data, reducing manual review effort

Manipal Institute of Technology

Research Intern

Manipal, KA, India

Oct. 2023 – Jan. 2024

- Engineered a 5-layer **CNN** (PyTorch, SCA-optimized) classifying lung CT scans at **92% accuracy**; published, winning **Best Paper Award** at ICDATA 2024 (London)

PROJECTS

CineMatch – Movie Recommender | Python, FastAPI, PostgreSQL, Docker, AWS | GitHub

July 2026 – Aug. 2026

- Deployed a hybrid recommender (**ALS** collaborative filtering blended with **TF-IDF** content similarity) as a Dockerized **FastAPI** service on **AWS EC2** with a **PostgreSQL** backend and a public Streamlit demo
- Validated personalization gains via a pre-registered **A/B experiment** on held-out data, measuring a **44.9%** relative lift in hit-rate@10 over a popularity baseline ($p < 0.001$)

IntelliTalent – Resume to Job Matching | Python, GPT-4o-mini, FAISS, AWS EC2 | GitHub

Feb. 2026 – Apr. 2026

- Developed a resume-to-job matching platform (**GPT-4o-mini**, **FAISS RAG**) ranking listings via hybrid cosine-similarity/skill-overlap scoring
- Replaced per-job LLM calls with regex taxonomy matching, cutting latency from **5+ seconds** to **~50ms** per job; deployed with Docker on **AWS EC2**

Social Media Hate Speech and Toxicity Mining | Python, NLP, PyTorch, Scikit-learn | GitHub

Jan. 2026 – April 2026

- Processed **1.8M+** comments (Jigsaw + Twitter) into a stratified **79K**-sample to fix severe class imbalance (**91.7%** non-toxic), then resolved a **2.6%** vocabulary-overlap fine-tuning failure (stalled at **67%** accuracy) by retraining **DistilBERT** end-to-end on the balanced corpus
- Outperformed classical baselines (best **ROC-AUC 0.915**, Logistic Regression) with the fine-tuned **DistilBERT**, achieving **F1 = 0.870** and **91%** recall on held-out data

Loan Approval Prediction | Python, Scikit-learn, SHAP, Pandas | GitHub

July 2026

- Tuned Decision Tree and Random Forest across **216** hyperparameter sets in an sklearn Pipeline for **83.7% accuracy** and **0.796 ROC AUC**; **SHAP** flagged Credit History as the dominant signal