

Himanshu Manish Jain

+1 (303) 641-6694 | himanshumjain15@gmail.com | [LinkedIn](#) | [Portfolio](#) | Boulder, CO | Open to Relocate

EDUCATION

University of Colorado Boulder

Master of Science in Data Science

- Relevant Courses: Statistical Inference, Data Mining, Big Data Architecture

Manipal Institute of Technology

Bachelor of Science in Electronics and Communication

Boulder, CO

Aug. 2025 – May 2027

India

Aug. 2021 – May 2025

SKILLS

Programming & Tools: Python, SQL, PostgreSQL, Docker, FastAPI, AWS, Tableau, Power BI, R

Data & ML: Pandas, NumPy, scikit-learn, PyTorch, CNN, NLP, LLM APIs

Statistical Methods: Statistical Modeling, Hypothesis Testing, A/B Testing, Anomaly Detection, Bootstrapping, Exploratory Data Analysis (EDA)

Certifications: Google Data Analytics Professional Certificate, Data Analysis With Tidiverse (CU Boulder)

EXPERIENCE

Summer Atlantic Capital

Data Science Intern

Newark, DE

June 2026 – Present

- Cut compensation lookup time from **10–20 minutes** to seconds, saving hundreds of analyst-hours across the People Analytics org, by building an end-to-end **Python/SQL** platform with a precomputed pass and two-tier caching layer over **617K** rows and **173K** external market records spanning **150** job families
- Delivered self-service pay-equity and workforce analytics to **10,000+** employees and leadership by designing executive **Tableau** dashboards, engineering a **MAD/modified-Z** outlier and privacy-suppression statistical engine, and integrating the **Claude API** for entity resolution with human-in-the-loop safeguards
- Recovered **~1,600** previously dropped market records and restored full data coverage by diagnosing and fixing two silent defects in the data pipeline

Ernst and Young (EY)

Data Science Intern

Bengaluru, India

Jan. 2025 – July 2025

- Validated a **22%** adoption lift by running **A/B tests** on **100K+** platform users to evaluate an auto call routing feature, using a **two-sample t-test** to confirm significance before scaling
- Increased adoption by **30%** across Tele-MANAS, a **national** mental health platform, by transforming **10M+** healthcare records into actionable insights via **SQL**, surfacing regional drop-off patterns, and orchestrating **cross-functional** rollouts with medical, engineering, and government stakeholders
- Reduced strategic decision cycles by **25%** for **10+** regional teams by designing and deploying **Power BI** dashboards to track performance KPIs

LG Electronics

Data Analyst Intern

Noida, India

May 2024 – June 2024

- Identified a **10%** cost-saving opportunity by applying **z-score outlier analysis** to flag suppliers with anomalous raw material costs, enabling data-driven negotiations
- Accelerated contract approvals by **15%** by designing a scalable **Python + SQL** pipeline to integrate and process supplier **KPI** data, reducing manual review effort

Manipal Institute of Technology

Research Intern

Manipal, India

Oct. 2023 – Jan. 2024

- Engineered a 5-layer **CNN** using PyTorch, optimized with a Sine Cosine Algorithm (SCA), classifying lung CT scans with **92% accuracy**; published at the 5th International Conference on Data Analytics (London, 2024), winning **Best Paper Award**

PROJECTS

CineMatch – Movie Recommender | Python, FastAPI, PostgreSQL, Docker, AWS | [GitHub](#)

July 2026 – Aug. 2026

- Built and deployed a hybrid recommender (**ALS** collaborative filtering blended with **TF-IDF** content similarity) as a Dockerized **FastAPI** service on **AWS EC2** with a **PostgreSQL** backend and a public Streamlit demo
- Validated personalization through a pre-registered **A/B experiment** on held-out data, measuring a **44.9%** relative lift in hit-rate@10 over a popularity baseline ($p < 0.001$)

IntelliTalent – Resume to Job Matching | Python, GPT-4o-mini, FAISS, AWS EC2 | [GitHub](#)

Feb. 2026 – Apr. 2026

- Built an end-to-end resume-to-job matching platform using GPT-4o-mini over a **FAISS-backed RAG** pipeline, ranking listings with a hybrid cosine similarity / skill-overlap score
- Replaced per-job LLM calls with regex taxonomy matching, cutting latency from **5+ seconds** to **~50ms** per job; deployed with Docker on AWS EC2

Social Media Hate Speech and Toxicity Mining | Python, NLP, Scikit-learn, Seaborn | [GitHub](#)

Jan. 2026 – April 2026

- Classified hate speech across 7 toxicity categories on **1.77M+** records via TF-IDF, SVD embeddings, and KMeans, surfacing **97.4%** vocabulary divergence and **91.5%** class imbalance
- Built an **NLP** pipeline distinguishing targeted hate speech from generic offensive language across **1.8M+** records (Jigsaw + TweetEval, **2.6%** vocabulary overlap), benchmarking classical baselines (Naive Bayes, Logistic Regression, SVM, LightGBM) on 200-dim **TF-IDF/LSA** features against a fine-tuned **DistilBERT** via precision, recall, F1, and confusion-matrix analysis

Loan Approval Prediction | Python, Scikit-learn, SHAP, Pandas | [GitHub](#)

July 2026

- Tuned Decision Tree and Random Forest across **216** hyperparameter sets in an sklearn Pipeline for **83.7% accuracy** and **0.796 ROC AUC**; **SHAP** flagged Credit History as the dominant signal